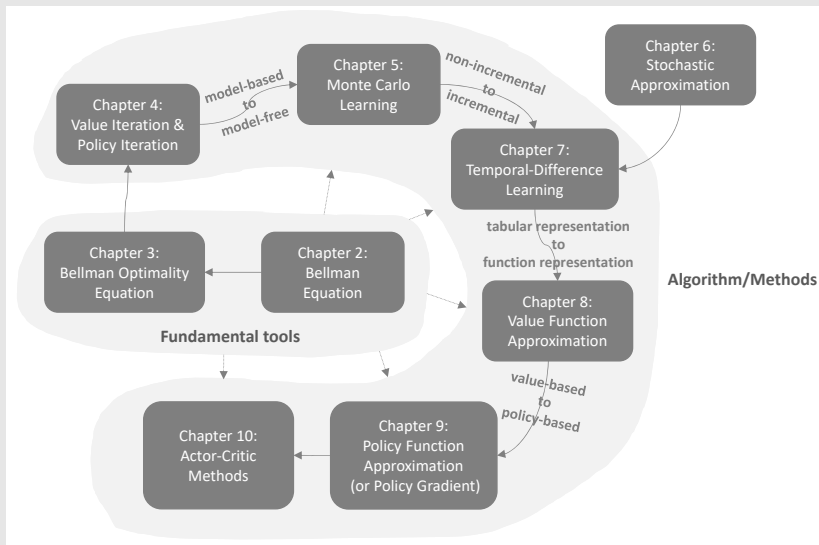


Lecture 10: Actor-Critic Methods

Shiyu Zhao

Introduction



Actor-critic methods are still policy gradient methods.

- They emphasize the structure that incorporates the policy gradient and value-based methods.

What are “actor” and “critic”?

- Here, “actor” refers to [policy update](#). It is called *actor* is because the policies will be applied to take actions.
- Here, “critic” refers to [policy evaluation](#) or [value estimation](#). It is called *critic* because it criticizes the policy by evaluating it.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The simplest actor-critic

Revisit the idea of policy gradient introduced in the last lecture.

- 1) A scalar metric $J(\theta)$, which can be \bar{v}_π or \bar{r}_π .
- 2) The gradient-ascent algorithm maximizing $J(\theta)$ is

$$\begin{aligned}\theta_{t+1} &= \theta_t + \alpha \nabla_\theta J(\theta_t) \\ &= \theta_t + \alpha \mathbb{E}_{S \sim \eta, A \sim \pi} \left[\nabla_\theta \ln \pi(A|S, \theta_t) q_\pi(S, A) \right]\end{aligned}$$

- 3) The stochastic gradient-ascent algorithm is

$$\theta_{t+1} = \theta_t + \alpha \nabla_\theta \ln \pi(a_t|s_t, \theta_t) q_t(s_t, a_t)$$

We can see “actor” and “critic” from this algorithm:

- This algorithm corresponds to actor!
- The algorithm estimating $q_t(s, a)$ corresponds to critic!

The simplest actor-critic

How to get $q_t(s_t, a_t)$?

So far, we have studied **two ways** to estimate action values:

- **Monte Carlo learning:** If MC is used, the corresponding algorithm is called **REINFORCE** or **Monte Carlo policy gradient**.
 - We introduced in the last lecture.
- **Temporal-difference learning:** If TD is used, such kind of algorithms are usually called **actor-critic**.
 - We will introduce in this lecture.

The simplest actor-critic

The simplest actor-critic algorithm (QAC)

Aim: Search for an optimal policy by maximizing $J(\theta)$.

At time step t in each episode, do

Generate a_t following $\pi(a|s_t, \theta_t)$, observe r_{t+1}, s_{t+1} , and then generate a_{t+1} following $\pi(a|s_{t+1}, \theta_t)$.

Critic (value update):

$$w_{t+1} = w_t + \alpha_w [r_{t+1} + \gamma q(s_{t+1}, a_{t+1}, w_t) - q(s_t, a_t, w_t)] \nabla_w q(s_t, a_t, w_t)$$

Actor (policy update):

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \ln \pi(a_t | s_t, \theta_t) q(s_t, a_t, w_{t+1})$$

The simplest actor-critic

Remarks:

- The critic corresponds to “SARSA+value function approximation”.
- The actor corresponds to the policy update algorithm.
- The algorithm is **on-policy** (why is PG on-policy?).
 - Since the policy is stochastic, no need to use techniques like ϵ -greedy.
- This particular actor-critic algorithm is sometimes referred to as **Q Actor-Critic** (QAC).
- Though simple, this algorithm reveals the core idea of actor-critic methods. It can be extended to generate many other algorithms as shown later.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Next, we extend QAC to advantage actor-critic (A2C)

- The **core idea** is to **introduce a baseline to reduce variance**.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Baseline invariance

Property: the policy gradient is invariant to an additional baseline:

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{S \sim \eta, A \sim \pi} \left[\nabla_{\theta} \ln \pi(A|S, \theta_t) q_{\pi}(S, A) \right] \\ &= \mathbb{E}_{S \sim \eta, A \sim \pi} \left[\nabla_{\theta} \ln \pi(A|S, \theta_t) (q_{\pi}(S, A) - b(S)) \right]\end{aligned}$$

Here, the additional baseline $b(S)$ is a scalar function of S .

Next, we answer two questions:

- Why is it valid?
- Why is it useful?

Baseline invariance

First, why is it valid?

That is because

$$\mathbb{E}_{S \sim \eta, A \sim \pi} \left[\nabla_{\theta} \ln \pi(A|S, \theta_t) b(S) \right] = 0$$

The details:

$$\begin{aligned} \mathbb{E}_{S \sim \eta, A \sim \pi} \left[\nabla_{\theta} \ln \pi(A|S, \theta_t) b(S) \right] &= \sum_{s \in \mathcal{S}} \eta(s) \sum_{a \in \mathcal{A}} \pi(a|s, \theta_t) \nabla_{\theta} \ln \pi(a|s, \theta_t) b(s) \\ &= \sum_{s \in \mathcal{S}} \eta(s) \sum_{a \in \mathcal{A}} \nabla_{\theta} \pi(a|s, \theta_t) b(s) \\ &= \sum_{s \in \mathcal{S}} \eta(s) b(s) \sum_{a \in \mathcal{A}} \nabla_{\theta} \pi(a|s, \theta_t) \\ &= \sum_{s \in \mathcal{S}} \eta(s) b(s) \nabla_{\theta} \sum_{a \in \mathcal{A}} \pi(a|s, \theta_t) \\ &= \sum_{s \in \mathcal{S}} \eta(s) b(s) \nabla_{\theta} 1 = 0 \end{aligned}$$

Second, why is the baseline useful?

The gradient is $\nabla_{\theta} J(\theta) = \mathbb{E}[X]$ where

$$X(S, A) \doteq \nabla_{\theta} \ln \pi(A|S, \theta_t) [q_{\pi}(S, A) - b(S)]$$

We have

- $\mathbb{E}[X]$ is invariant to $b(S)$.
- $\text{var}(X)$ is NOT invariant to $b(S)$.
 - Why? Because $\text{tr}[\text{var}(X)] = \mathbb{E}[X^T X] - \bar{x}^T \bar{x}$ and

$$\begin{aligned}\mathbb{E}[X^T X] &= \mathbb{E} [(\nabla_{\theta} \ln \pi)^T (\nabla_{\theta} \ln \pi) (q_{\pi}(S, A) - b(S))^2] \\ &= \mathbb{E} [\|\nabla_{\theta} \ln \pi\|^2 (q_{\pi}(S, A) - b(S))^2]\end{aligned}$$

Imagine b is huge (e.g., 1 million)

Our goal: Select an **optimal baseline** b to minimize $\text{var}(X)$

- **Benefit:** when we use a random sample to approximate $\mathbb{E}[X]$, the estimation variance would also be small.

In the algorithms of REINFORCE and QAC,

- There is no baseline.
- Or, we can say $b = 0$, **which is not guaranteed to be a good baseline.**

Baseline invariance

- The **optimal baseline** that can minimize $\text{var}(X)$ is, for any $s \in \mathcal{S}$,

$$b^*(s) = \frac{\mathbb{E}_{A \sim \pi} [\|\nabla_{\theta} \ln \pi(A|s, \theta_t)\|^2 q_{\pi}(s, A)]}{\mathbb{E}_{A \sim \pi} [\|\nabla_{\theta} \ln \pi(A|s, \theta_t)\|^2]}.$$

See the proof in my book.

- Although this baseline is optimal, it is complex.
- We can remove the weight $\|\nabla_{\theta} \ln \pi(A|s, \theta_t)\|^2$ and select the suboptimal baseline:

$$b(s) = \mathbb{E}_{A \sim \pi} [q_{\pi}(s, A)] = v_{\pi}(s)$$

which is the state value of s !

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The algorithm of advantage actor-critic

When $b(s) = v_\pi(s)$,

- the gradient-ascent algorithm is

$$\begin{aligned}\theta_{t+1} &= \theta_t + \alpha \mathbb{E} \left[\nabla_\theta \ln \pi(A|S, \theta_t) [q_\pi(S, A) - v_\pi(S)] \right] \\ &\doteq \theta_t + \alpha \mathbb{E} \left[\nabla_\theta \ln \pi(A|S, \theta_t) \delta_\pi(S, A) \right]\end{aligned}$$

where

$$\delta_\pi(S, A) \doteq q_\pi(S, A) - v_\pi(S)$$

is called the **advantage function** (why called advantage?).

- the stochastic version of this algorithm is

$$\begin{aligned}\theta_{t+1} &= \theta_t + \alpha \nabla_\theta \ln \pi(a_t|s_t, \theta_t) [q_t(s_t, a_t) - v_t(s_t)] \\ &= \theta_t + \alpha \nabla_\theta \ln \pi(a_t|s_t, \theta_t) \delta_t(s_t, a_t)\end{aligned}$$

The algorithm of advantage actor-critic

Moreover, the algorithm can be reexpressed as

$$\begin{aligned}\theta_{t+1} &= \theta_t + \alpha \nabla_{\theta} \ln \pi(a_t | s_t, \theta_t) \delta_t(s_t, a_t) \\ &= \theta_t + \alpha \frac{\nabla_{\theta} \pi(a_t | s_t, \theta_t)}{\pi(a_t | s_t, \theta_t)} \delta_t(s_t, a_t) \\ &= \theta_t + \underbrace{\alpha \left(\frac{\delta_t(s_t, a_t)}{\pi(a_t | s_t, \theta_t)} \right)}_{\text{step size}} \nabla_{\theta} \pi(a_t | s_t, \theta_t)\end{aligned}$$

- The step size is proportional to the **relative value** δ_t rather than the **absolute value** q_t , which is more reasonable.
- It can still **well balance exploration and exploitation**.

The algorithm of advantage actor-critic

Furthermore, the advantage function is approximated by the TD error:

$$\delta_t = q_t(s_t, a_t) - v_t(s_t) \rightarrow r_{t+1} + \gamma v_t(s_{t+1}) - v_t(s_t)$$

- This approximation is **reasonable** because

$$\mathbb{E}[q_\pi(S, A) - v_\pi(S) | S = s_t, A = a_t] = \mathbb{E}\left[R + \gamma v_\pi(S') - v_\pi(S) | S = s_t, A = a_t\right]$$

- **Benefit:** only need one network to approximate $v_\pi(s)$ rather than two networks for $q_\pi(s, a)$ and $v_\pi(s)$.

The algorithm of advantage actor-critic

Advantage actor-critic (A2C) or TD actor-critic

Aim: Search for an optimal policy by maximizing $J(\theta)$.

At time step t in each episode, do

Generate a_t following $\pi(a|s_t, \theta_t)$ and then observe r_{t+1}, s_{t+1} .

TD error (advantage function):

$$\delta_t = r_{t+1} + \gamma v(s_{t+1}, w_t) - v(s_t, w_t)$$

Critic (value update):

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w v(s_t, w_t)$$

Actor (policy update):

$$\theta_{t+1} = \theta_t + \alpha_\theta \delta_t \nabla_\theta \ln \pi(a_t | s_t, \theta_t)$$

It is on-policy. Since the policy $\pi(\theta_t)$ is stochastic, no need to use techniques like ε -greedy.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Introduction

- Policy gradient is on-policy.
 - Why? because the gradient is $\nabla_{\theta} J(\theta) = \mathbb{E}_{S \sim \eta, A \sim \pi}[*]$
- Can we convert it to off-policy?
 - Yes, by **importance sampling**
 - The importance sampling technique is not limited to AC, but also to any algorithm that aims to estimate an expectation.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Illustrative examples

Consider a random variable $X \in \mathcal{X} = \{+1, -1\}$.

If the probability distribution of X is p_0 :

$$p_0(X = +1) = 0.5, \quad p_0(X = -1) = 0.5$$

then the expectation of X is

$$\mathbb{E}_{X \sim p_0}[X] = (+1) \cdot 0.5 + (-1) \cdot 0.5 = 0.$$

Question: how to estimate $\mathbb{E}[X]$ by using some samples $\{x_i\}$?

Illustrative examples

Case 1 (we are already familiar):

- The samples $\{x_i\}$ are generated according to p_0 :

$$\mathbb{E}[x_i] = \mathbb{E}[X], \quad \text{var}[x_i] = \text{var}[X]$$

Then, the average value can converge to the expectation:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \mathbb{E}[X], \quad \text{as } n \rightarrow \infty$$


because

$$\mathbb{E}[\bar{x}] = \mathbb{E}[X], \quad \text{var}[\bar{x}] = \frac{1}{n} \text{var}[X]$$

See my book for details (Law of large numbers).

Illustrative examples

Figure: Samples and $\bar{x} \rightarrow \mathbb{E}[X]$



Illustrative examples

Case 2 (a new case that we want to study):

- The samples $\{x_i\}$ are generated according to **another distribution** p_1 :

$$p_1(X = +1) = 0.8, \quad p_1(X = -1) = 0.2$$

The expectation is

$$\mathbb{E}_{X \sim p_1}[X] = (+1) \cdot 0.8 + (-1) \cdot 0.2 = 0.6$$

If we use the average of the samples, then without suprising

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \mathbb{E}_{X \sim p_1}[X] = 0.6 \neq \mathbb{E}_{X \sim p_0}[X]$$

Illustrative examples

Question: Can we use $\{x_i\} \sim p_1$ to estimate $\mathbb{E}_{X \sim p_0}[X]$?

- **Why to do that?**

We may want to estimate $\mathbb{E}_{A \sim \pi}[*]$ where π is the *target policy* based on the samples of a *behavior policy* β .

- **How to do that?**


- We can't achieve that if directly using \bar{x} :

$$\bar{x} \rightarrow \mathbb{E}_{X \sim p_1}[X] = 0.6 \neq \mathbb{E}_{X \sim p_0}[X]$$

- We can achieve that by using the importance sampling technique.

Illustrative examples

Figure: Samples and $\bar{x} \rightarrow \mathbb{E}_{X \sim p_1}[X]$ (the dotted line)



- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic**
 - Illustrative examples
 - Importance sampling**
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Importance sampling

Note that

$$\mathbb{E}_{X \sim p_0}[X] = \sum_x p_0(x)x = \sum_x p_1(x) \underbrace{\frac{p_0(x)}{p_1(x)}}_{f(x)} x = \mathbb{E}_{X \sim p_1}[f(X)]$$

- Thus, we can estimate $\mathbb{E}_{X \sim p_1}[f(X)]$ in order to estimate $\mathbb{E}_{X \sim p_0}[X]$.
- How to estimate $\mathbb{E}_{X \sim p_1}[f(X)]$? Easy. Let

$$\bar{f} \doteq \frac{1}{n} \sum_{i=1}^n f(x_i), \quad \text{where } x_i \sim p_1$$

Then,

$$\begin{aligned}\mathbb{E}_{X \sim p_1}[\bar{f}] &= \mathbb{E}_{X \sim p_1}[f(X)] \\ \text{var}_{X \sim p_1}[\bar{f}] &= \frac{1}{n} \text{var}_{X \sim p_1}[f(X)]\end{aligned}$$

Importance sampling

Therefore, \bar{f} is a good approximation for $\mathbb{E}_{X \sim p_1}[f(X)] = \mathbb{E}_{X \sim p_0}[f(X)]$

$$\mathbb{E}_{X \sim p_0}[f(X)] \approx \bar{f} = \frac{1}{n} \sum_{i=1}^n f(x_i) = \frac{1}{n} \sum_{i=1}^n \frac{p_0(x_i)}{p_1(x_i)} f(x_i)$$

- $\frac{p_0(x_i)}{p_1(x_i)}$ is called the *importance weight*.
 - If $p_1(x_i) = p_0(x_i)$, the importance weight is one and \bar{f} becomes \bar{x} .
 - If $p_0(x_i) \geq p_1(x_i)$, x_i can be more often sampled by p_0 than p_1 .
The importance weight (> 1) can emphasize the importance of this sample.

Importance sampling

You may ask: While $\bar{f} = \frac{1}{n} \sum_{i=1}^n \frac{p_0(x_i)}{p_1(x_i)} x_i$ requires $p_0(x)$, if I know $p_0(x)$, why not directly calculate the expectation?

Answer: It is applicable to the case where it is easy to calculate $p_0(x)$ given an x , but difficult to calculate the expectation.

- For example, continuous case, complex expression of p_0 , or no expression of p_0 (e.g., p_0 represented by a neural network).

Importance sampling

Summary: if $\{x_i\} \sim p_1$,

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \mathbb{E}_{X \sim p_1}[X]$$

$$\bar{f} = \frac{1}{n} \sum_{i=1}^n \frac{p_0(x_i)}{p_1(x_i)} x_i \rightarrow \mathbb{E}_{X \sim p_0}[X]$$




Figure: Blue dotted line: \bar{x} ; Red solid line: \bar{f}

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic**
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient**
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The theorem of off-policy policy gradient

Like the previous on-policy case, we need to derive the policy gradient in the off-policy case.

- Suppose β is the behavior policy that generates experience samples.
- Our aim is to use these samples to update a target policy π that can minimize the metric

$$J(\theta) = \sum_{s \in \mathcal{S}} d_{\beta}(s) v_{\pi}(s) = \mathbb{E}_{S \sim d_{\beta}} [v_{\pi}(S)]$$

where d_{β} is the stationary distribution under policy β .

The theorem of off-policy policy gradient

Theorem (Off-policy policy gradient theorem)

In the discounted case where $\gamma \in (0, 1)$, the gradient of $J(\theta)$ is

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{S \sim \rho, A \sim \beta} \left[\frac{\pi(A|S, \theta)}{\beta(A|S)} \nabla_{\theta} \ln \pi(A|S, \theta) q_{\pi}(S, A) \right]$$

where β is the behavior policy and ρ is a state distribution.

See the details and the proof in my book.

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic**
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic**
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The algorithm of off-policy actor-critic

The off-policy policy gradient is also **invariant to a baseline $b(s)$** .

- In particular, we have

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{S \sim \rho, A \sim \beta} \left[\frac{\pi(A|S, \theta)}{\beta(A|S)} \nabla_{\theta} \ln \pi(A|S, \theta) (q_{\pi}(S, A) - b(S)) \right]$$

- To reduce the estimation variance, we can select the baseline as $b(S) = v_{\pi}(S)$ and obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[\frac{\pi(A|S, \theta)}{\beta(A|S)} \nabla_{\theta} \ln \pi(A|S, \theta) (q_{\pi}(S, A) - v_{\pi}(S)) \right]$$

The algorithm of off-policy actor-critic

The corresponding stochastic gradient-ascent algorithm is

$$\theta_{t+1} = \theta_t + \alpha_\theta \frac{\pi(a_t|s_t, \theta_t)}{\beta(a_t|s_t)} \nabla_\theta \ln \pi(a_t|s_t, \theta_t) (q_t(s_t, a_t) - v_t(s_t))$$

Similar to the on-policy case,

$$q_t(s_t, a_t) - v_t(s_t) \approx r_{t+1} + \gamma v_t(s_{t+1}) - v_t(s_t) \doteq \delta_t(s_t, a_t)$$

Then, the algorithm becomes

$$\theta_{t+1} = \theta_t + \alpha_\theta \frac{\pi(a_t|s_t, \theta_t)}{\beta(a_t|s_t)} \nabla_\theta \ln \pi(a_t|s_t, \theta_t) \delta_t(s_t, a_t)$$

and hence

$$\theta_{t+1} = \theta_t + \alpha_\theta \left(\frac{\delta_t(s_t, a_t)}{\beta(a_t|s_t)} \right) \nabla_\theta \pi(a_t|s_t, \theta_t)$$

The algorithm of off-policy actor-critic

Off-policy actor-critic based on importance sampling

Initialization: A given behavior policy $\beta(a|s)$. A target policy $\pi(a|s, \theta_0)$ where θ_0 is the initial parameter vector. A value function $v(s, w_0)$ where w_0 is the initial parameter vector.

Aim: Search for an optimal policy by maximizing $J(\theta)$.

At time step t in each episode, do

Generate a_t following $\beta(s_t)$ and then observe r_{t+1}, s_{t+1} .

TD error (advantage function):

$$\delta_t = r_{t+1} + \gamma v(s_{t+1}, w_t) - v(s_t, w_t)$$

Critic (value update):

$$w_{t+1} = w_t + \alpha_w \frac{\pi(a_t|s_t, \theta_t)}{\beta(a_t|s_t)} \delta_t \nabla_w v(s_t, w_t)$$

Actor (policy update):

$$\theta_{t+1} = \theta_t + \alpha_\theta \frac{\pi(a_t|s_t, \theta_t)}{\beta(a_t|s_t)} \delta_t \nabla_\theta \ln \pi(a_t|s_t, \theta_t)$$

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

Up to now, the policies used in the policy gradient methods are all **stochastic** since $\pi(a|s, \theta) > 0$ for every (s, a) .

Can we use **deterministic policies** in the policy gradient methods?

- Benefit: it can handle continuous action.

Introduction

The ways to represent a policy:

- Up to now, a general policy is denoted as $\pi(a|s, \theta) \in [0, 1]$, which can be either stochastic or deterministic.
- Now, the deterministic policy is specifically denoted as

$$a = \mu(s, \theta) \doteq \mu(s)$$

- μ is a mapping from \mathcal{S} to \mathcal{A} .
- μ can be represented by, for example, a neural network with the input as s , the output as a , and the parameter as θ .
- We may write $\mu(s, \theta)$ in short as $\mu(s)$.

Outline

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The theorem of deterministic policy gradient

- The policy gradient theorems introduced before are merely valid for stochastic policies.
- If the policy must be deterministic, we must derive a new policy gradient theorem.
- The ideas and procedures are similar.

The theorem of deterministic policy gradient

Consider the metric of average state value in the discounted case:

$$J(\theta) = \mathbb{E}[v_\mu(s)] = \sum_{s \in \mathcal{S}} d_0(s) v_\mu(s)$$

where $d_0(s)$ is a probability distribution satisfying $\sum_{s \in \mathcal{S}} d_0(s) = 1$.

- d_0 is selected to be independent of μ . The gradient in this case is easier to calculate.
- There are two special yet important cases of selecting d_0 .
 - The first special case is that $d_0(s_0) = 1$ and $d_0(s \neq s_0) = 0$, where s_0 is a specific starting state of interest.
 - The second special case is that d_0 is the stationary distribution of a behavior policy that is different from the μ .

The theorem of deterministic policy gradient

Theorem (Deterministic policy gradient theorem in the discounted case)

In the discounted case where $\gamma \in (0, 1)$, the gradient of $J(\theta)$ is

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \sum_{s \in \mathcal{S}} \rho_{\mu}(s) \nabla_{\theta} \mu(s) (\nabla_a q_{\mu}(s, a))|_{a=\mu(s)} \\ &= \mathbb{E}_{S \sim \rho_{\mu}} [\nabla_{\theta} \mu(S) (\nabla_a q_{\mu}(S, a))|_{a=\mu(S)}]\end{aligned}$$

Here, ρ_{μ} is a state distribution.

See more details and the proof in my book.

One important difference from the stochastic case:

- The gradient does not involve the distribution of the action A (why?).
- As a result, the deterministic policy gradient method is **off-policy**.

Outline

- 1 The simplest actor-critic (QAC)
- 2 Advantage actor-critic (A2C)
 - Baseline invariance
 - The algorithm of advantage actor-critic
- 3 Off-policy actor-critic
 - Illustrative examples
 - Importance sampling
 - The theorem of off-policy policy gradient
 - The algorithm of off-policy actor-critic
- 4 Deterministic actor-critic (DPG)
 - The theorem of deterministic policy gradient
 - The algorithm of deterministic actor-critic

The algorithm of deterministic actor-critic

Based on the policy gradient, the gradient-ascent algorithm for maximizing $J(\theta)$ is:

$$\theta_{t+1} = \theta_t + \alpha_{\theta} \mathbb{E}_{S \sim \rho_{\mu}} [\nabla_{\theta} \mu(S) (\nabla_a q_{\mu}(S, a)) |_{a=\mu(S)}]$$

The corresponding stochastic gradient-ascent algorithm is

$$\theta_{t+1} = \theta_t + \alpha_{\theta} \nabla_{\theta} \mu(s_t) (\nabla_a q_{\mu}(s_t, a)) |_{a=\mu(s_t)}$$

The algorithm of deterministic actor-critic

Deterministic actor-critic algorithm

Initialization: A given behavior policy $\beta(a|s)$. A deterministic target policy $\mu(s, \theta_0)$ where θ_0 is the initial parameter vector. A value function $v(s, w_0)$ where w_0 is the initial parameter vector.

Aim: Search for an optimal policy by maximizing $J(\theta)$.

At time step t in each episode, do

Generate a_t following β and then observe r_{t+1}, s_{t+1} .

TD error:

$$\delta_t = r_{t+1} + \gamma q(s_{t+1}, \mu(s_{t+1}, \theta_t), w_t) - q(s_t, a_t, w_t)$$

Critic (value update):

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w q(s_t, a_t, w_t)$$

Actor (policy update):

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu(s_t, \theta_t) (\nabla_a q(s_t, a, w_{t+1}))|_{a=\mu(s_t)}$$

The algorithm of deterministic actor-critic

Remarks:

- This is an off-policy implementation where the behavior policy β may be different from μ .
- β can also be replaced by $\mu + \text{noise}$.
- How to select the function to represent $q(s, a, w)$?
 - **Linear function**: $q(s, a, w) = \phi^T(s, a)w$ where $\phi(s, a)$ is the feature vector. Details can be found in the DPG paper.
 - **Neural networks**: deep deterministic policy gradient (DDPG) method.

Summary

- The simplest actor-critic
- Advantage actor-critic
- Off-policy actor-critic
- Deterministic actor-critic

The end

This is the end of the course, but a start for your journey
in the field of reinforcement learning!